

Kokkos Tutorial

Christian R. Trott ¹, Dan Ibanez ¹, David S. Hollman ¹, Dan Sunderland ¹, Duane Labreche ¹, Nathan Ellingwood ¹, Steve Bova ¹, Graham Lopez ², Galen Shipman ³, and Geoffrey Womeldorff ³

¹Sandia National Laboratories, ²Oak Ridge National Laboratory

³Los Alamos National Laboratory

Oak Ridge National Laboratories July 23-27, 2018

Sandia National Laboratories is a multi-mission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC., a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.

SAND2019-xxxx

Task parallelism

Fine-grained dependent execution.

Learning objectives:

- ▶ Basic interface for fine-grained tasking in Kokkos
- ▶ How to express dynamic dependency structures in Kokkos tasking
- ▶ When to use Kokkos tasking

Recall that **data parallel** code is composed of a **pattern**, a **policy**, and a **functor**

```
Kokkos::parallel_for(  
  Kokkos::RangePolicy<>(exec_space, 0, N),  
  SomeFunctor()  
);
```

Task parallel code similarly has a **pattern**, a **policy**, and a **functor**

```
Kokkos::task_spawn(  
  Kokkos::TaskSingle(scheduler, TaskPriority::High),  
  SomeFunctor()  
);
```

```
struct MyTask {  
    using value_type = double;  
    template <class TeamMember>  
    KOKKOS_INLINE_FUNCTION  
    void operator()(TeamMember& member, double& result);  
};
```

- ▶ Tell Kokkos what the **value type** of your task's output is.
- ▶ Take a **team member** argument, analogous to the team member passed in by `Kokkos::TeamPolicy` in hierarchical parallelism
- ▶ The **output** is expressed by assigning to a parameter, similar to with `Kokkos::parallel_reduce`

- ▶ `Kokkos::TaskSingle()`
 - ▶ Run the task with a single worker thread
- ▶ `Kokkos::TaskTeam()`
 - ▶ Run the task with all of the threads in a team
 - ▶ Think of it like being inside of a `parallel_for` with a `TeamPolicy`
- ▶ Both policies take a scheduler, an optional predecessor, and an optional priority (more on schedulers and predecessors later)

- ▶ `Kokkos::task_spawn()`
 - ▶ `Kokkos::host_spawn()` (same thing, but from host memory space)
 - ▶ Soon, we'll have just `scheduler.spawn()`
- ▶ `Kokkos::respawn()`
 - ▶ **Argument order is backwards; policy comes second!**
 - ▶ Soon, we'll have just `scheduler.respawn()`
- ▶ `task_spawn()` and `host_spawn()` return a `Kokkos::Future` representing the completion of the task (see next slide), which can be used as a predecessor to another operation.

How do futures and dependencies work?

```
struct MyTask {
    using value_type = double;
    Kokkos::Future<double, Kokkos::DefaultExecutionSpace> dep;
    int depth;
    KOKKOS_INLINE_FUNCTION MyTask(int d) : depth(d) { }
    template <class TeamMember>
    KOKKOS_INLINE_FUNCTION
    void operator()(TeamMember& member, double& result) {
        if(depth == 1) result = 3.14;
        else if(dep.is_null()) {
            dep =
                Kokkos::task_spawn(
                    Kokkos::TaskSingle(member.scheduler()),
                    MyTask(depth-1)
                );
            Kokkos::respawn(*this, dep);
        }
        else {
            result = depth * dep.get();
        }
    }
};
```

```
template <class Scheduler>
struct MyTask {
    using value_type = double;
    Kokkos::BasicFuture<double, Scheduler> dep;
    int depth;
    KOKKOS_INLINE_FUNCTION MyTask(int d) : depth(d) { }
    template <class TeamMember>
    KOKKOS_INLINE_FUNCTION
    void operator()(TeamMember& member, double& result);
};
```

Available Schedulers:

- ▶ TaskScheduler<ExecSpace>
- ▶ TaskSchedulerMultiple<ExecSpace>
- ▶ ChaseLevTaskScheduler<ExecSpace>


```
using execution_space = Kokkos::DefaultExecutionSpace;
using scheduler_type = Kokkos::TaskScheduler<execution_space>;
using memory_space = scheduler_type::memory_space;
using memory_pool_type = scheduler_type::memory_pool;
size_t memory_pool_size = 1 << 22;

auto scheduler =
    scheduler_type(memory_pool_type(memory_pool_size));

Kokkos::BasicFuture<double, scheduler_type> result =
    Kokkos::host_spawn(
        Kokkos::TaskSingle(scheduler),
        MyTask<scheduler_type>(10)
    );
Kokkos::wait(scheduler);
printf("Result is %f", result.get());
```

- ▶ Tasks always run to completion
- ▶ There is no way to wait or block inside of a task
 - ▶ `future.get()` does not block!
- ▶ Tasks that do not `respawn` themselves are complete
 - ▶ The value in the `result` parameter is made available through `future.get()` to any dependent tasks.
- ▶ The second argument to `respawn` can only be either a predecessor (future) or a scheduler, not a proper execution policy
 - ▶ We are fixing this to provide a more consistent overload in the next release.
- ▶ Tasks can only have one predecessor (at a time)
 - ▶ Use `scheduler.when_all()` to aggregate predecessors (see next slide)

```
using void_future =  
    Kokkos::BasicFuture<void, scheduler_type>;  
auto f1 =  
    Kokkos::task_spawn(Kokkos::TaskSingle(scheduler), X{});  
auto f2 =  
    Kokkos::task_spawn(Kokkos::TaskSingle(scheduler), Y{});  
void_future f_array[] = { f1, f2 };  
void_future f_12 = scheduler.when_all(f_array, 2);  
auto f3 =  
    Kokkos::task_spawn(  
        Kokkos::TaskSingle(scheduler, f_12), FuncXY{}  
    );
```

- ▶ To create an aggregate Future, use `scheduler.when_all()`
- ▶ `scheduler.when_all()` always returns a void future.
- ▶ (Also, any future is implicitly convertible to a void future of the same Scheduler type)

Formula

$$F_N = F_{N-1} + F_{N-2}$$

$$F_0 = 0$$

$$F_1 = 1$$

Serial algorithm

```
int fib(int n) {  
    if(n < 2) return n;  
    else {  
        return fib(n-1) + fib(n-2);  
    }  
}
```

Details:

- ▶ Location: Intro-Full/Exercises/08
- ▶ Implement the FibonacciTask task functor recursively
- ▶ Spawn the root task from the host and wait for the scheduler to make it ready

Hints:

- ▶ Do the F_{N-1} and F_{N-2} subproblems in separate tasks
- ▶ Use a `scheduler.when_all()` to wait on the subproblems